Mario Lavina

GitHub LinkedIn

Education

Imperial College London — **London, UK**
MSc. Artificial Intelligence — 2021–2022

London School of Economics — **London, UK**
BSc. Mathematics with Economics — 2018–2021

Experience

FINESSE — **Los Angeles, CA**
Engineering Manager, ML Engineer — Jan 2023–Present

Lead engineering at a $300M Series A startup (20 headcount, $16M revenue). Streamlined team from 8 to 5 while doubling release cadence. Built CI/CD, code review, and sprint processes from scratch.

- Built two internal dashboards: a **Next.js** operations view (launched September 2026) joining our design platform, logistics systems and **Stripe Payouts**, so staff see each creator design, its estimated China-to-US arrival and payout in one place, and a monitor for every agent and automation I built.
- Built AI-powered QA system using Claude agents with Figma design verification and Metro simulator — reviews **40 PRs/week**, replaced 2 QA engineers (**$160K/yr**), cut release cycles in half, reduced production bugs by **83%**.
- Built demand forecasting (XGBoost + SageMaker) that triggers reorders across 24K SKUs and 30 suppliers, cutting overproduction **30%** and saving **$1M/yr**; planning the weekly China-to-US shipment from it moved 35% of volume from air to sea, cutting freight cost per unit **21%** with stockouts flat.
- Shipped mobile app in 3 months after inheriting a stalled project (React Native), grew it from $0 to **$300K/mo revenue** — **10K WAU**, 6× higher conversion than web (4.2% vs 0.7%), now driving 23% of company revenue at 16% higher AOV. Built and tested its **TypeScript** API (Express, TypeORM, Postgres) with Testcontainers integration tests in CI.
- Shipped a text-to-image model (**Qwen-Image**) serving **40K generations/day** at 4s latency under 60 concurrent requests, adopted by 72% of WAU; users who created and posted AI-generated designs showed **2× D30 retention**.
- Built the labeling tool for a **1M-image** training set: a model pre-fills each label and labelers confirm or fix it with keyboard shortcuts, cutting time per image from **40s to 15s** while agreement on a 2,000-image double-labeled check held at **94%**, saving **6,900 labeler hours**.
- Trained a caption model on **CLIP** image features with 8,000 stylist-written captions; in a blind test on 500 fashion images, stylists preferred its captions over GPT-4V's **64%** of the time, and its tags filled missing color, fabric and fit fields on **40,000 product pages**.
- Built a Ray pipeline that ingests **100K+ TikTok and Instagram posts a day**, extracts category, purchase intent and sentiment from their comments, and feeds a real-time engine ranking fashion trends by likes, comments, sentiment and author, with an average rank error of **5** in validation.
- Ran a 2-day, no-downtime AWS-to-Google Cloud move that cut the cloud bill from **$10K to $1.7K a month**: deleted unused resources, tuned the database onto a smaller server and archived old data.

Imperial Driverless — **London, UK**
Lead ML Engineer — Sep 2021–Jul 2022

- Led a 5-engineer Motion Planning team to a **top-5 finish** in Imperial Driverless's first-ever competition, with **80% run completion**; owned path-and-motion planning on a Jetson Xavier race car at up to 30 mph.
- Built a **real-time SLAM** system (ORB-SLAM2 + occupancy mapping), hitting sub-50ms pose latency and <5cm drift / km in GPS-denied environments.
- Implemented LiDAR-camera sensor-fusion module (**ICP + EKF**, Velodyne VLP-16 + camera), cutting track-boundary localization error **4x (0.5m →0.12m)** and tightening the racing-line planning tolerance.
- Shipped a fused perception stack - (**YOLOv3**) cone detection + (**MiDaS**) monocular depth - sustaining 37 FPS on Jetson within the car's power budget, feed the planner at control-loop rate.

Research

Generating Safety Critical Scenarios, *Dr Dario Paccagnan and Dr Panagiotis Angeloudis*

– Developed RL framework (MADDPG, DQN) synthesizing edge-case accident scenarios for autonomous vehicles, achieving **7% higher collision coverage** than existing scenario generators. Distributed training across 16×A100 GPUs using Ray RLlib with **4× faster convergence**.

Selected Projects

Auto Parts Finder, *VibeCon 2026 hackathon (175 of 2,000 applicants selected)* Oct 2026

– Built an **AI voice agent** (ElevenLabs) that phones car dealerships for mechanics: from a VIN and part, it gets price and stock in a **median 2.5 min**, against 9 min for 4 LA mechanics (60 timed calls, 1 in 3 unanswered). A second voice agent answers dealers' return calls, matched to the request by caller number.

– In **30 live calls** to 12 dealers it got **17 quotes, 16 at the exact price** (checked against the recordings), and a repair shop ran 14 real requests through it in a day, getting 9 quotes. Tuned on 50 simulated edge cases (43/50 to **50/50**), it scored **7/8** on 8 of the 20 real dealer calls, held out from tuning.

Skills

Programming Languages: Python, SQL, Mojo (experimenting)
Frameworks: PyTorch, Tensorflow, OpenCV, NLTK, HuggingFace (transformers, diffusers), RLlib, OpenAI Gym, Ray
Software/Tools: AWS Stack (S3, EC2, RDS, Sagemaker, Cloudfront, Quicksight), Pinecone, W&B, FastAPI, Render